



Career Change Prediction: EDA & Insights

Introduction

This dataset explores factors influencing career changes, including academic background, job satisfaction, industry growth, and more. With over 30,000 records and 22 attributes, the target variable, Likely to Change Occupation, predicts whether an individual is likely to switch careers.

This analysis aims to uncover patterns and provide actionable insights into career transitions.

Import Library

```
In [3]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import seaborn as sns
```

Uploading Csv file

```
In [4]: df = pd.read_csv(r"C:\Users\Syed Arif\OneDrive\Desktop\Carrer Change\career_ch
ange_prediction_dataset.csv")
```

Data Preprocessing

head()

head is used show to the By default = 5 rows in the dataset

```
In [5]: df.head()
```

Out[5]:

	Field of Study	Current Occupation	Age	Gender	Years of Experience	Education Level	Industry Growth Rate	Job Satisfaction	Work-Life Balance
0	Medicine	Business Analyst	48	Male	7	High School	High	7	10
1	Education	Economist	44	Male	26	Master's	Low	10	3
2	Education	Biologist	21	Female	27	Master's	Low	8	3
3	Education	Business Analyst	33	Male	14	PhD	Medium	7	9
4	Arts	Doctor	28	Female	0	PhD	Low	3	1

5 rows × 23 columns

.tail()

tail is used to show rows by Descending order

```
In [6]: df.tail()
```

Out[6]:

	Field of Study	Current Occupation	Age	Gender	Years of Experience	Education Level	Industry Growth Rate	Job Satisfaction	Work-Life Balance
38439	Biology	Business Analyst	24	Female	34	High School	Low	8	
38440	Mechanical Engineering	Artist	21	Female	24	High School	Low	2	
38441	Computer Science	Mechanical Engineer	35	Female	21	High School	High	4	
38442	Arts	Business Analyst	35	Male	11	PhD	Medium	9	
38443	Law	Mechanical Engineer	37	Male	23	Master's	Medium	6	

5 rows × 23 columns

shape

```
In [7]: df.shape
```

```
Out[7]: (38444, 23)
```

Columns

```
In [8]: df.columns
```

```
Out[8]: Index(['Field of Study', 'Current Occupation', 'Age', 'Gender',  
              'Years of Experience', 'Education Level', 'Industry Growth Rate',  
              'Job Satisfaction', 'Work-Life Balance', 'Job Opportunities', 'Salar  
y',  
              'Job Security', 'Career Change Interest', 'Skills Gap',  
              'Family Influence', 'Mentorship Available', 'Certifications',  
              'Freelancing Experience', 'Geographic Mobility',  
              'Professional Networks', 'Career Change Events', 'Technology Adoptio  
n',  
              'Likely to Change Occupation'],  
             dtype='object')
```

.dtypes

This Attribute show the data type of each column

```
In [9]: df.dtypes
```

```
Out[9]: Field of Study          object
Current Occupation          object
Age                        int64
Gender                    object
Years of Experience        int64
Education Level            object
Industry Growth Rate       object
Job Satisfaction           int64
Work-Life Balance          int64
Job Opportunities          int64
Salary                    int64
Job Security               int64
Career Change Interest     int64
Skills Gap                 int64
Family Influence           object
Mentorship Available       int64
Certifications             int64
Freelancing Experience      int64
Geographic Mobility         int64
Professional Networks      int64
Career Change Events       int64
Technology Adoption        int64
Likely to Change Occupation int64
dtype: object
```

.unique()

In a column, It show the unique value of specific column.

```
In [10]: df["Current Occupation"].unique()
```

```
Out[10]: array(['Business Analyst', 'Economist', 'Biologist', 'Doctor', 'Lawyer',
                'Software Developer', 'Artist', 'Psychologist', 'Teacher',
                'Mechanical Engineer'], dtype=object)
```

.nunique()

It will show the total no of unique value from whole data frame

.value_counts

It Shows all the unique values with their count

```
In [11]: df["Current Occupation"].value_counts()

Out[11]: Current Occupation
Software Developer      3892
Psychologist            3890
Doctor                  3888
Teacher                 3886
Artist                  3881
Business Analyst        3858
Mechanical Engineer     3827
Lawyer                  3781
Biologist               3774
Economist               3767
Name: count, dtype: int64
```

isnull()

It shows the how many null values

```
In [12]: df.isnull()

Out[12]:
```

	Field of Study	Current Occupation	Age	Gender	Years of Experience	Education Level	Industry Growth Rate	Job Satisfaction	Work-Life Balance
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...
38439	False	False	False	False	False	False	False	False	False
38440	False	False	False	False	False	False	False	False	False
38441	False	False	False	False	False	False	False	False	False
38442	False	False	False	False	False	False	False	False	False
38443	False	False	False	False	False	False	False	False	False

38444 rows × 23 columns

.info()

To Show Data type of each column

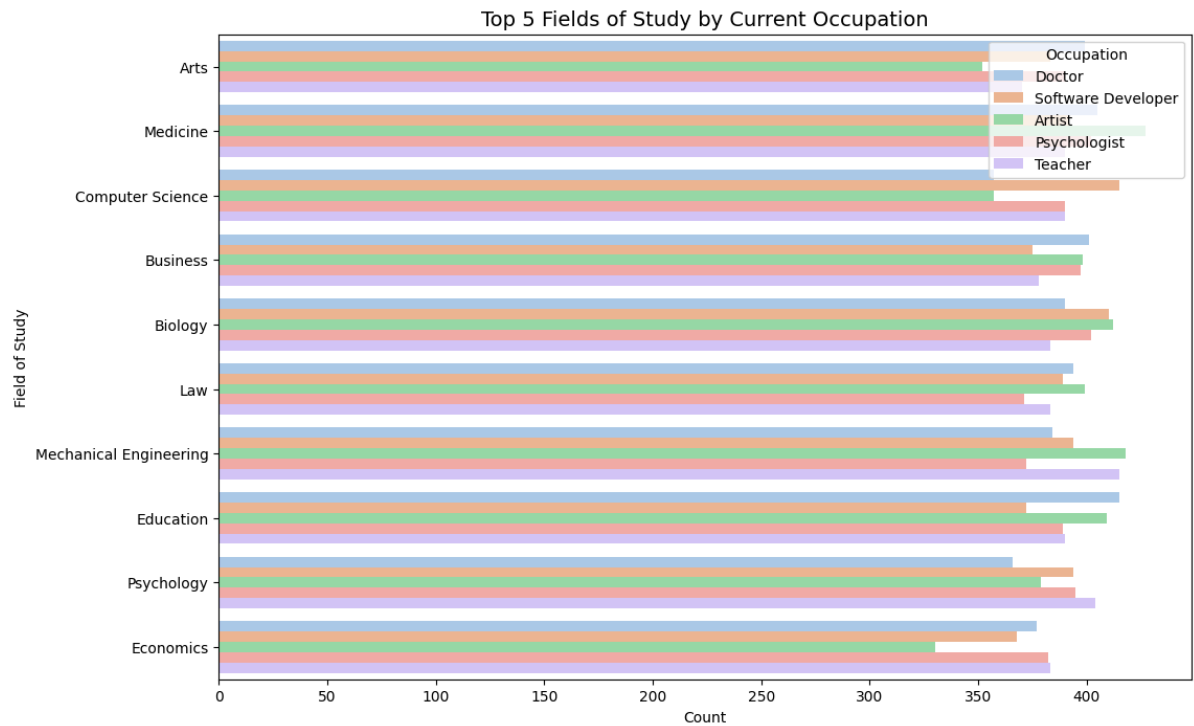
In [13]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38444 entries, 0 to 38443
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Field of Study                        38444 non-null  object
1   Current Occupation                    38444 non-null  object
2   Age                                  38444 non-null  int64
3   Gender                               38444 non-null  object
4   Years of Experience                   38444 non-null  int64
5   Education Level                       38444 non-null  object
6   Industry Growth Rate                  38444 non-null  object
7   Job Satisfaction                      38444 non-null  int64
8   Work-Life Balance                    38444 non-null  int64
9   Job Opportunities                    38444 non-null  int64
10  Salary                               38444 non-null  int64
11  Job Security                          38444 non-null  int64
12  Career Change Interest                38444 non-null  int64
13  Skills Gap                            38444 non-null  int64
14  Family Influence                      28812 non-null  object
15  Mentorship Available                  38444 non-null  int64
16  Certifications                        38444 non-null  int64
17  Freelancing Experience                 38444 non-null  int64
18  Geographic Mobility                   38444 non-null  int64
19  Professional Networks                 38444 non-null  int64
20  Career Change Events                  38444 non-null  int64
21  Technology Adoption                   38444 non-null  int64
22  Likely to Change Occupation            38444 non-null  int64
dtypes: int64(17), object(6)
memory usage: 6.7+ MB
```

1. Field of Study and Current Occupation (Top 5):

```
In [14]: top_occupations = df["Current Occupation"].value_counts().head(5).index
top_fields = df[df["Current Occupation"].isin(top_occupations)]

plt.figure(figsize=(12, 8))
sns.countplot(data=top_fields, y="Field of Study", hue="Current Occupation", palette="pastel")
plt.title("Top 5 Fields of Study by Current Occupation", fontsize=14)
plt.xlabel("Count")
plt.ylabel("Field of Study")
plt.legend(title="Occupation")
plt.show()
```

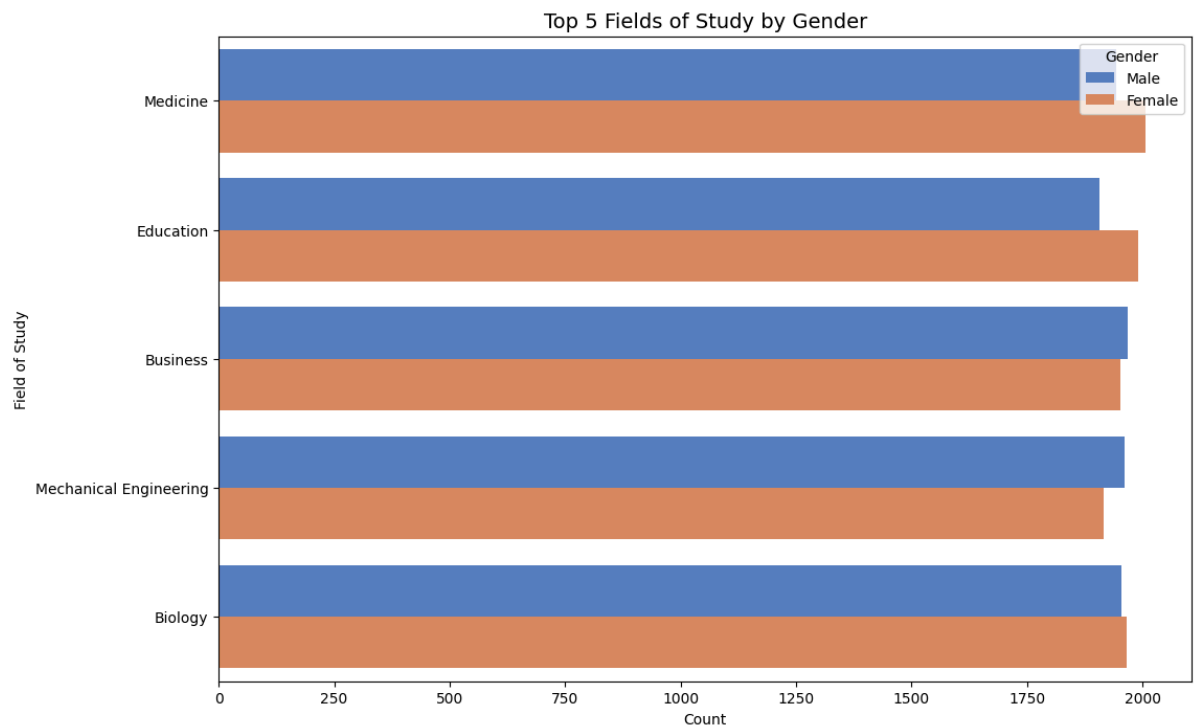


2. Field of Study and Gender Distribution (Top 5 Fields):

Question: Field of Study and Gender


```
In [18]: top_field_study = df["Field of Study"].value_counts().head(5).index
field_gender = df[df["Field of Study"].isin(top_field_study)]

plt.figure(figsize=(12, 8))
sns.countplot(data=field_gender, y="Field of Study", hue="Gender", palette="muted")
plt.title("Top 5 Fields of Study by Gender", fontsize=14)
plt.xlabel("Count")
plt.ylabel("Field of Study")
plt.legend(title="Gender")
plt.show()
```



3. Occupation and Gender Distribution (Top 5 Occupations):

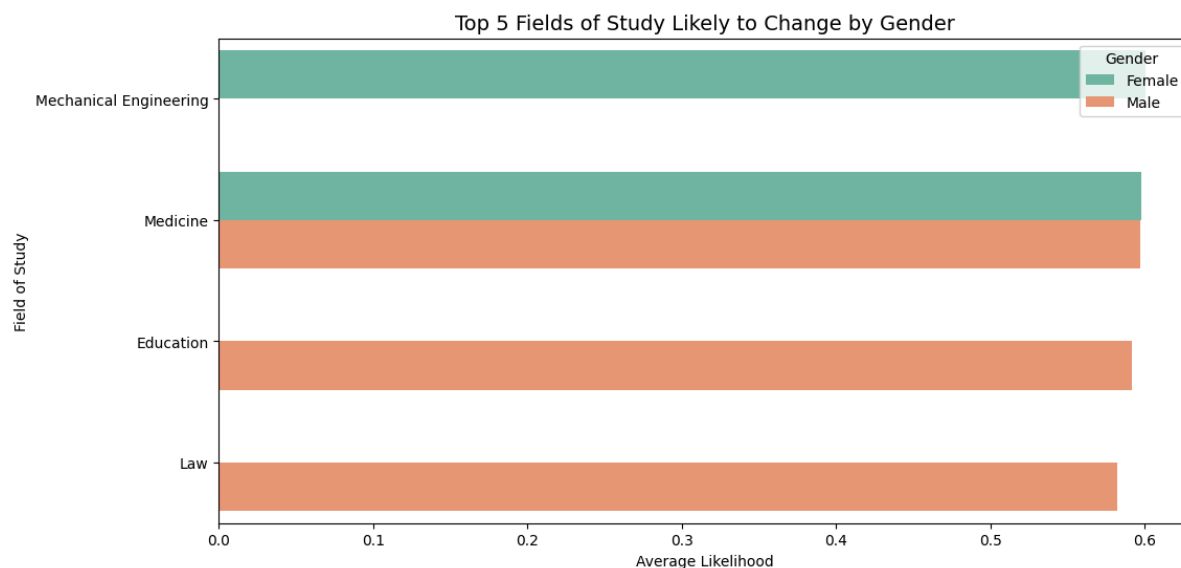

```
In [40]: plt.figure(figsize=(12, 8))
sns.countplot(data=top_fields, y="Current Occupation", hue="Gender", palette="coolwarm")
plt.title("Top 5 Occupations by Gender", fontsize=14)
plt.xlabel("Count")
plt.ylabel("Occupation")
plt.legend(title="Gender")
plt.show()
```



4. Top 5 Fields of Study Changed by Occupation and Gender:

```
In [20]: study_change = df.groupby(["Field of Study", "Gender"])["Likely to Change Occupation"].mean().nlargest(5).reset_index()

plt.figure(figsize=(12, 6))
sns.barplot(data=study_change, x="Likely to Change Occupation", y="Field of Study", hue="Gender", palette="Set2")
plt.title("Top 5 Fields of Study Likely to Change by Gender", fontsize=14)
plt.xlabel("Average Likelihood")
plt.ylabel("Field of Study")
plt.legend(title="Gender")
plt.show()
```



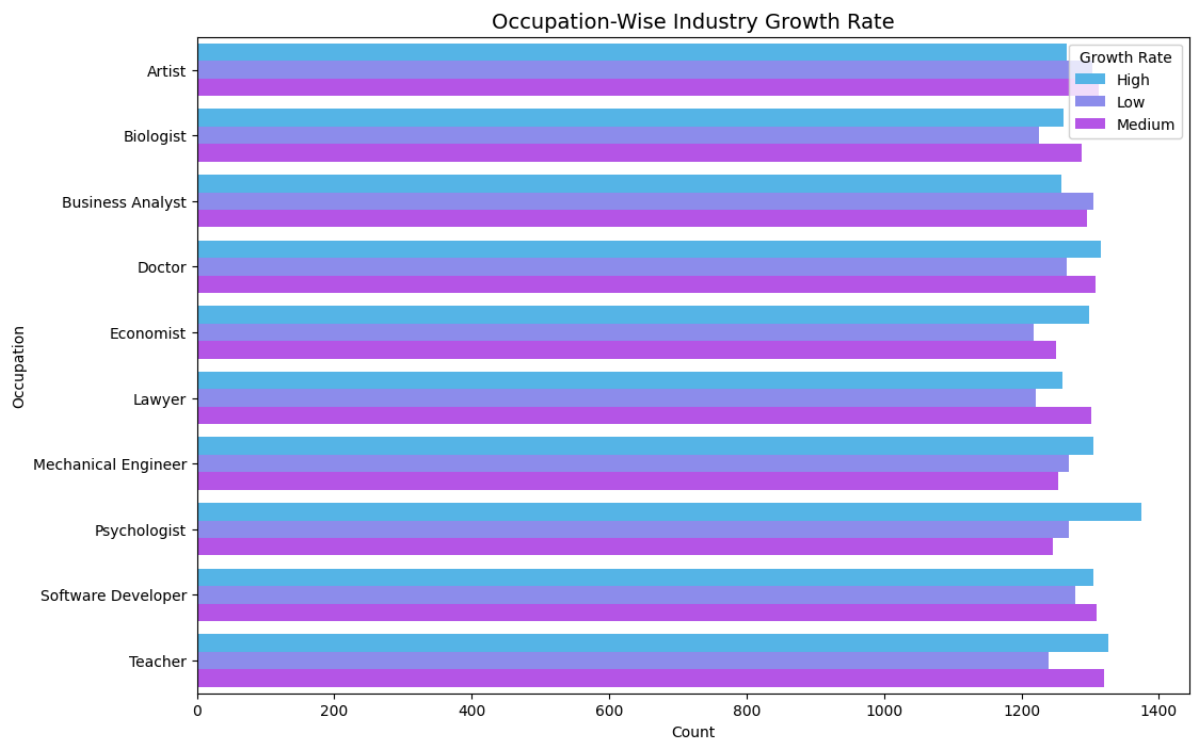
5. Occupation-Wise Industry Growth Rate:

```
In [39]: import matplotlib.pyplot as plt
import seaborn as sns

# Prepare the data for plotting
occupation_growth_rate = df.groupby(["Current Occupation", "Industry Growth Rate"]).size().reset_index(name="Count")

# Create a bar plot to show the count of each occupation by industry growth rate
plt.figure(figsize=(12, 8))
sns.barplot(data=occupation_growth_rate, x="Count", y="Current Occupation", hue="Industry Growth Rate", palette="cool")

# Customize the plot
plt.title("Occupation-Wise Industry Growth Rate", fontsize=14)
plt.xlabel("Count")
plt.ylabel("Occupation")
plt.legend(title="Growth Rate")
plt.show()
```



6. Gender-Wise Job Satisfaction:

Question: gender and calculate average job satisfaction?

```
In [38]: import matplotlib.pyplot as plt
import seaborn as sns

# Create a bar plot to show the average job satisfaction by gender
plt.figure(figsize=(10, 6))
sns.barplot(x="Gender", y="Job Satisfaction", data=df, palette="coolwarm", ci=
None)

# Customize the plot
plt.title("Gender-Wise Job Satisfaction", fontsize=14)
plt.xlabel("Gender")
plt.ylabel("Average Job Satisfaction")
plt.show()
```

C:\Users\Syed Arif\AppData\Local\Temp\ipykernel_7860\1016224037.py:6: FutureWarning:

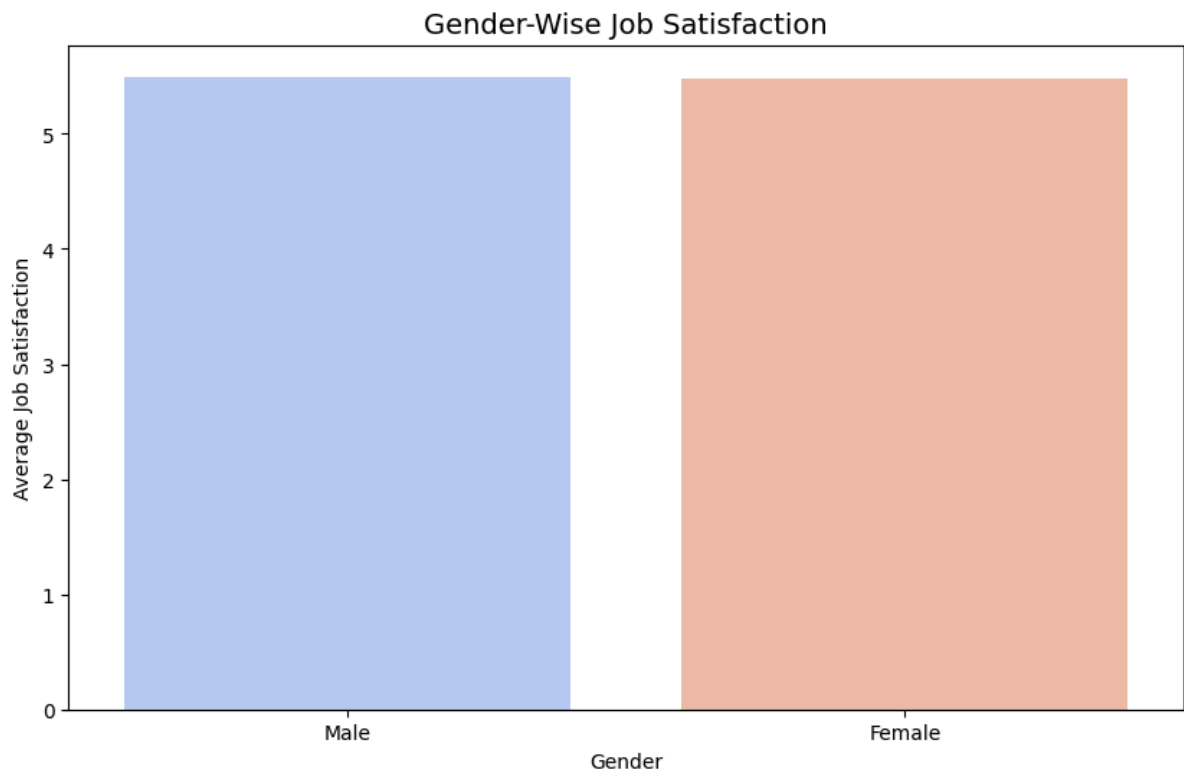
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x="Gender", y="Job Satisfaction", data=df, palette="coolwarm",
ci=None)
```

C:\Users\Syed Arif\AppData\Local\Temp\ipykernel_7860\1016224037.py:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x="Gender", y="Job Satisfaction", data=df, palette="coolwarm",
ci=None)
```



7. Occupation-Wise Salary

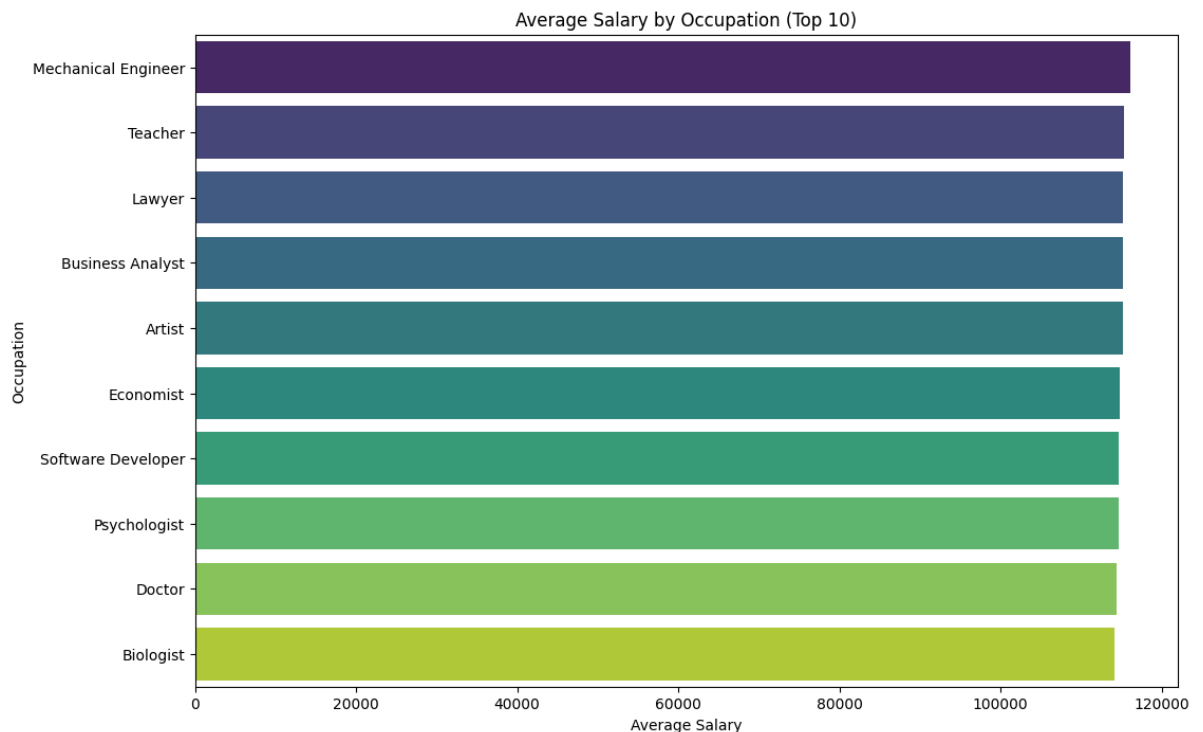
```
In [27]: occupation_salary_avg = df.groupby("Current Occupation")["Salary"].mean().sort
_values(ascending=False).head(10)

plt.figure(figsize=(12, 8))
sns.barplot(
    x=occupation_salary_avg.values,
    y=occupation_salary_avg.index,
    palette="viridis",
)
plt.title("Average Salary by Occupation (Top 10)")
plt.xlabel("Average Salary")
plt.ylabel("Occupation")
plt.show()
```

C:\Users\Syed Arif\AppData\Local\Temp\ipykernel_7860\2479832085.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



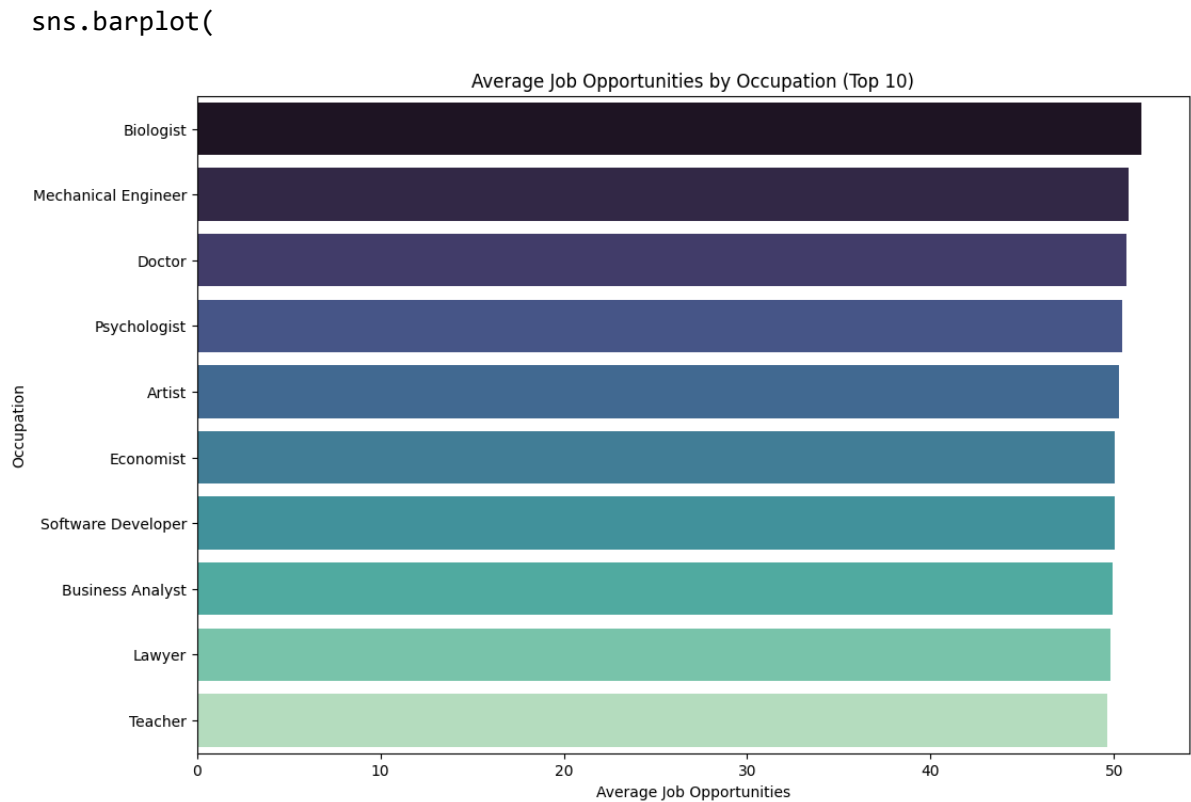
8. Occupation-Wise Job Opportunities:

```
In [29]: occupation_opportunities_avg = df.groupby("Current Occupation")["Job Opportunities"].mean().sort_values(ascending=False).head(10)

plt.figure(figsize=(12, 8))
sns.barplot(
    x=occupation_opportunities_avg.values,
    y=occupation_opportunities_avg.index,
    palette="mako",
)
plt.title("Average Job Opportunities by Occupation (Top 10)")
plt.xlabel("Average Job Opportunities")
plt.ylabel("Occupation")
plt.show()
```

C:\Users\Syed Arif\AppData\Local\Temp\ipykernel_7860\582844519.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.



9. Occupation-Wise Job Satisfaction:

```
In [36]: # Define a function to map numeric scores to satisfaction levels
def map_satisfaction(value):
    if value <= 3:
        return "Low"
    elif 4 <= value <= 7:
        return "Medium"
    else:
        return "High"

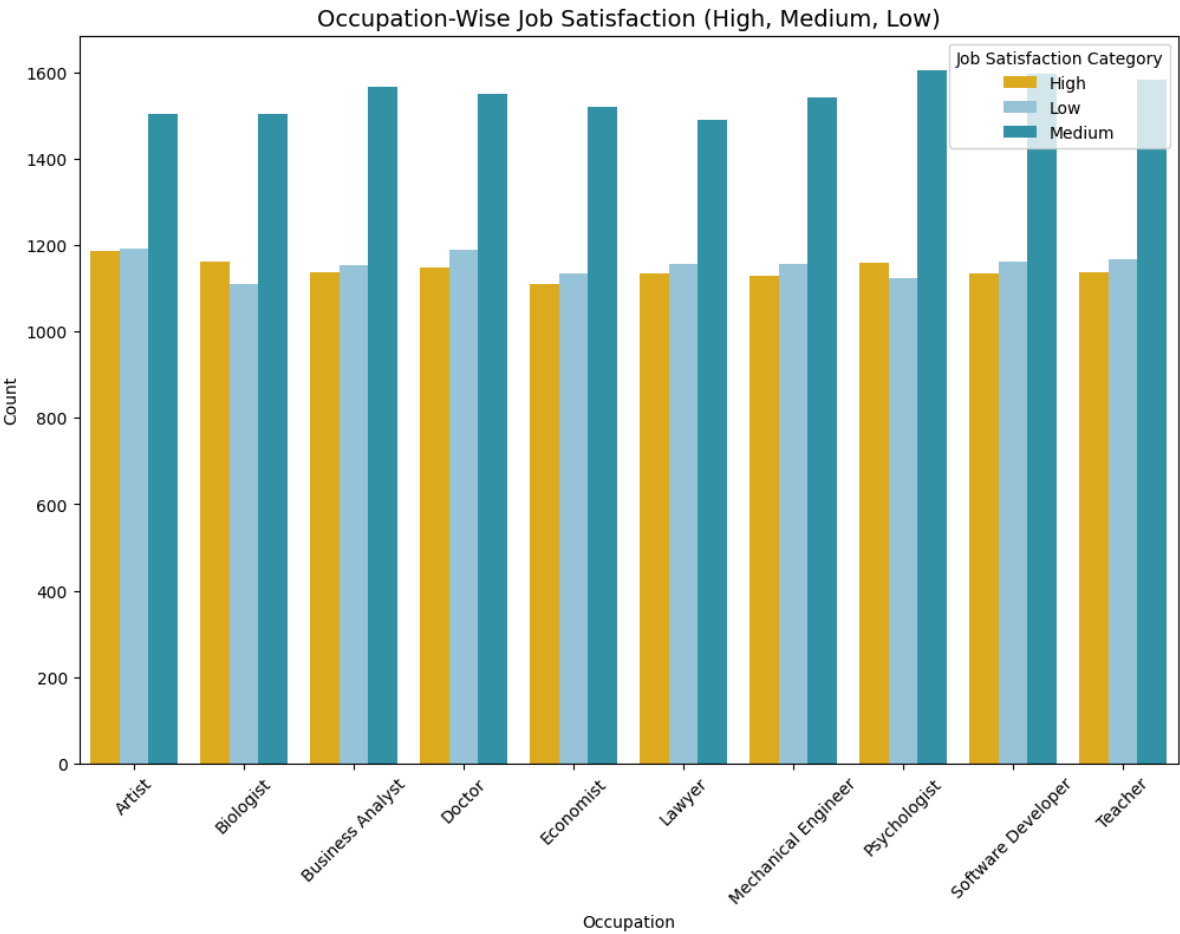
# Filter occupations ranked from 1 to 10 (e.g., top 10 by frequency)
top_occupations = df["Current Occupation"].value_counts().head(10).index
filtered_data = df[df["Current Occupation"].isin(top_occupations)]

# Map numeric satisfaction levels to categories
filtered_data["Job Satisfaction Category"] = filtered_data["Job Satisfaction"].apply(map_satisfaction)

# Prepare the data for visualization
occupation_satisfaction_counts = (
    filtered_data.groupby(["Current Occupation", "Job Satisfaction Category"])
    .size()
    .reset_index(name="Count")
)

# Plot a stacked bar chart using hue (similar to the first code example)
plt.figure(figsize=(12, 8))
sns.barplot(
    data=occupation_satisfaction_counts,
    x="Current Occupation",
    y="Count",
    hue="Job Satisfaction Category",
    palette=["#ffb703", "#8ecae6", "#219ebc"]
)

# Customize the plot
plt.title("Occupation-Wise Job Satisfaction (High, Medium, Low)", fontsize=14)
plt.xlabel("Occupation")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.legend(title="Job Satisfaction Category")
plt.show()
```

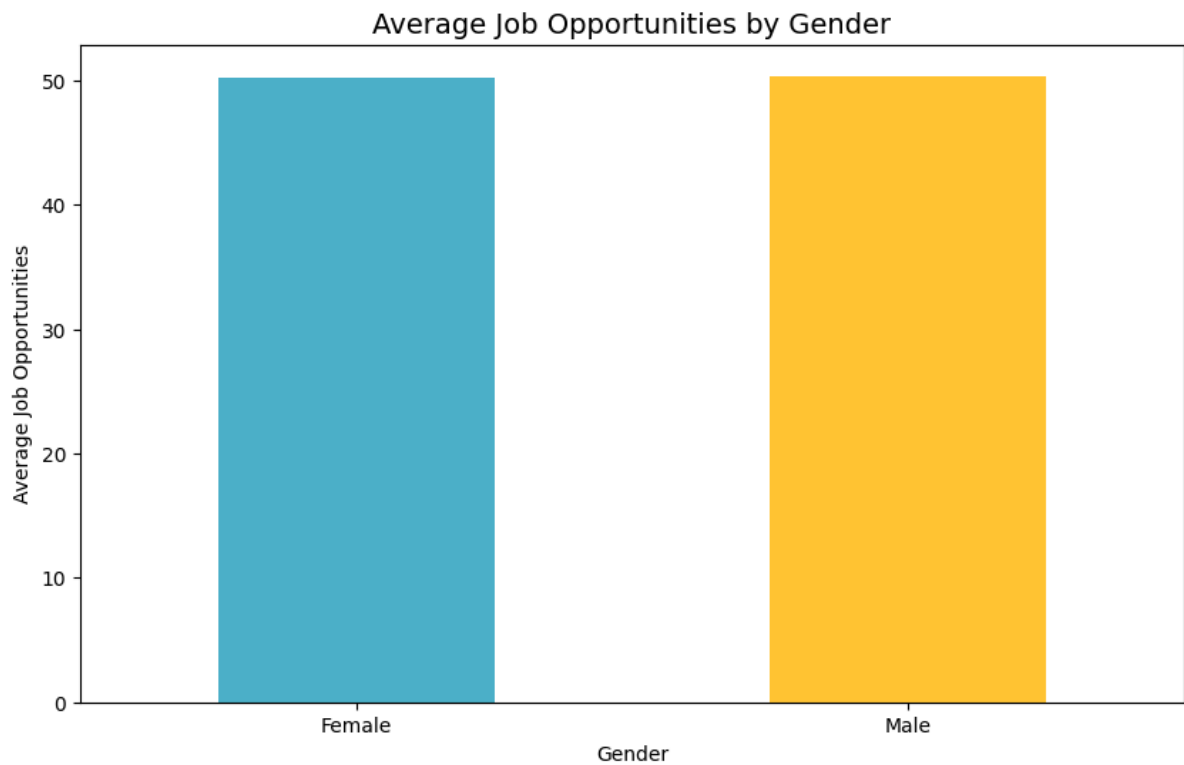


10. Occupation-Wise Industry Growth Rate

Question: How does job satisfaction vary across occupations and genders?


```
In [32]: gender_opportunities_mean = df.groupby("Gender")["Job Opportunities"].mean()

plt.figure(figsize=(10, 6))
gender_opportunities_mean.plot(
    kind="bar",
    color=["#219ebc", "#ffb703"],
    alpha=0.8,
)
plt.title("Average Job Opportunities by Gender", fontsize=14)
plt.xlabel("Gender")
plt.ylabel("Average Job Opportunities")
plt.xticks(rotation=0)
plt.show()
```



```
In [51]: # Import necessary Libraries
from wordcloud import WordCloud
import matplotlib.pyplot as plt

# Generate a word cloud based on the frequency of 'Current Occupation'
wordcloud = WordCloud(width=1000, height=600, background_color='white', colormap='viridis').generate_from_frequencies(df['Current Occupation'].value_counts())

# Plot the word cloud
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.title('Word Cloud of Occupations', fontsize=16)
plt.axis('off') # Disable the axis for better visualization
plt.show()
```

